



# Reinforcement Learning

a practical approach

Prof. Renzo Silva, Konstantin Kuchenmeister, Elijah Hall

*Department of Mathematics*  
*Columbia University*

Mathematics Building  
MATH203

Saturday 9<sup>th</sup> April, 2022

# Outline

- 1 Recap
- 2 Introduction to Reinforcement Learning
- 3 Q-learning
- 4 Deep Q-learning
- 5 Practical example: DQN Trading Bot

# Supervised Learning

## General Model

Recall that, in supervised learning we have a set up tuples

$$T = \{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\} \quad (1)$$

and through training we are trying to estimate the underlying map:

$$f(X) + \epsilon \longrightarrow Y \quad (2)$$

with

$$\hat{Y} = \hat{f}(X) \quad (3)$$

where  $\epsilon$  is the irreducible error we can not model.

# Supervised Learning

## Algorithms

To model this relationship we have, among others, looked at

### 1. Regression

- ▶ Linear Regression (Performance, Simplicity)
- ▶ kNN-Regressor (Performance, Simplicity)
- ▶ Decision Tree (Performance, Simplicity)
- ▶ Random Forest Regressor (Accuracy)
- ▶ Neural Network (Accuracy)

### 2. Classification

- ▶ Logistic Regression (Performance, Simplicity)
- ▶ kNN
- ▶ Decision Tree, Random Forest
- ▶ Neural Network (Accuracy)
- ▶ SVM

# Transitioning

BUT, in HW3 we have also looked at clustering...

What is the difference?

# Definition

## Reinforcement Learning

### Definition

Reinforcement learning problems involve learning what to do, how to map situations to actions, so as to maximize a numerical reward signal. [...] Moreover, the learner is not told which actions to take, as in many forms of machine learning, but instead must discover which actions yield the most reward by trying them out<sup>1</sup>.

In short: We also do not know the  $y'$ s... [► Motivation](#)

# Markov Decision Process

## Definitions

### Definition

In our context, a Markov Decision Process (MDP) is a 4-tuple  $(S, A, P, R)$  which allows us to formalize and model sequential decision making<sup>2</sup> as used in RL.

### Key Terminology

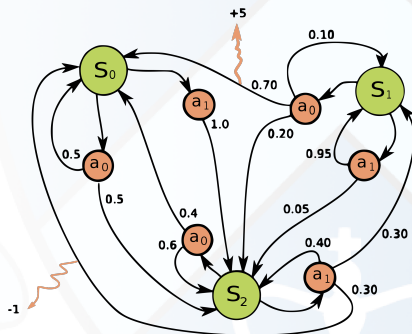
In particular,

- ▶  $S_i$  = State of the Environment
- ▶  $A_i$  = Set of actions available in state  $S_i$
- ▶  $P_i = \Pr(s_{t+1} = s' \mid s_t = s, a_t = a)$
- ▶  $R_i$  = Reward attained from transitioning from state  $S_i$  to  $S_{i+1}$  under action  $A_i$

# Markov Decision Process

Visualized

Figure: Markov Decision Process (Source=Wikipedia)





# Markov Decision Process

## Episodes

### Definition

Naturally, as we go through states through actions, we create a trajectory (similar to the natural filtration in stochastic processes):

$$\tau = (s_0, a_0, s_1, a_1, \dots) \quad (4)$$

If  $|\tau| < \infty$  (i.e.  $\exists$  a terminal state  $T$ ) it is called episode, and is used to describe the agent "playing the game" one time.

# Markov Decision Process

## Rewards

At every discrete time step we earn a reward:

### Lemma

$$R(s) = \mathbf{E} [R(t+1) \mid S_t = s'] \quad (5)$$

The infinite-horizon discounted return function:

$$G(t) = R(t+1) + \gamma R(t+2) + \dots = \sum_{k=0}^{\infty} \gamma^k R(t+k+1) \quad (6)$$

where  $\gamma$  is the discount factor.

# Markov Decision Process

## Formulating the problem

Now, what is the optimal action to take at an arbitrary time step  $t$ ?

$$\mathbf{E} [G(t) \mid S = s'] \quad (7)$$

Which corresponds to maximizing the sum of expected future rewards after starting from state  $s$ .

# Markov Decision Process

## Formulating the problem - Intro to the Bellman Equation

We can split this up in:

- ▶ immediate reward earned ( $R(t+1)$ )
- ▶ future rewards earned

$$\begin{aligned}v(s) &= \mathbf{E} [R(t+1) + \gamma R(t+2) + \gamma^2 R(t+3) + \dots \mid S = s'] \\&= \mathbf{E} [R(t+1) + \gamma(R(t+2) + \gamma R(t+3) + \dots) \mid S = s'] \\&= \mathbf{E} [R(t+1) + \gamma(G(t+1)) \mid S = s'] \\&= \mathbf{E} [R(t+1) + \gamma(v(S_{t+1})) \mid S = s']\end{aligned}$$

Which corresponds to maximizing the sum of expected future rewards after starting from state  $s$ . You enter a state, collect the immediate reward and average all the possible future rewards discounted.

# Markov Decision Process

## Policies

### Definition

A policy  $\pi$  is a distribution over actions given states,

$$\pi(a \mid s) = \Pr(A_t = a \mid S_t = s) \quad (8)$$

"If were in state  $s$ , the agent is going to choose action  $a$  with probability  $\pi(a \mid s)$ ." Defines the behavior of the agent.

# Markov Decision Process

## Policies

### Definition

The state-value function  $v_\pi(s)$  is the expected total return from starting from state  $s$  and following policy  $\pi$  thereafter:

$$v_\pi(s) = \mathbf{E}_\pi [G(t) \mid S_t = s] \quad (9)$$

### Definition

The action-value function  $q_\pi(s)$  is the expected total return from starting from state  $s$ , taking action  $a$  and following policy  $\pi$  thereafter:

$$q_\pi(s, a) = \mathbf{E}_\pi [G(t) \mid S_t = s, A_t = a] \quad (10)$$

# Markov Decision Process

## Policies

And again, we can decompose them into:

Lemma

Theorem

$$v_{\pi}(s) = \mathbf{E}_{\pi} [R(t+1) + \gamma(v(S_{t+1}) \mid S_t = s)]$$

Lemma

Theorem

$$q_{\pi}(s, a) = \mathbf{E}_{\pi} [R(t+1) + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a]$$

# Markov Decision Process

Relation between  $v$  and  $q$

## Remark

$$v_{\pi}(s) = \sum_{a \in A} \pi(a \mid s) q_{\pi}(s, a) \quad (11)$$

$$q_{\pi}(s) = R_s^a + \gamma \sum_{s \in S} P_{ss}^a \sum_{a \in A} \pi(a \mid s) q_{\pi}(s, a) \quad (12)$$



# Markov Decision Process

## Putting it all together - Optimality 1

So what is now the optimal way to act?

### Definition

$$v_*(s) = \max_{\pi} v_{\pi}(s) \quad (13)$$

$$q_*(s) = \max_{\pi} q_{\pi}(s, a) \quad (14)$$

The optimal is the maximum value of the respective function over all possible policies.

# Markov Decision Process

## Putting it all together - Optimality 2

### Definition

For any MDP, a policy  $\pi > \pi'$  iff  $v_\pi(s) > v_{\pi'}(s)$  and all optimal policies fulfill the previously defined optimal state-value  $v_*(s)$  and action-value  $q_*(s)$  functions.

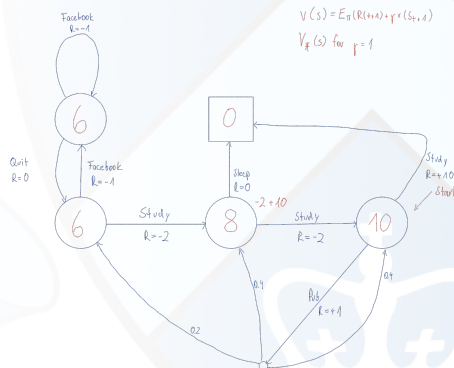
### Definition

$$\pi_* = 1_{\max_{a \in A} q_*(a,s)} \quad (15)$$

# Markov Decision Process

Visualized

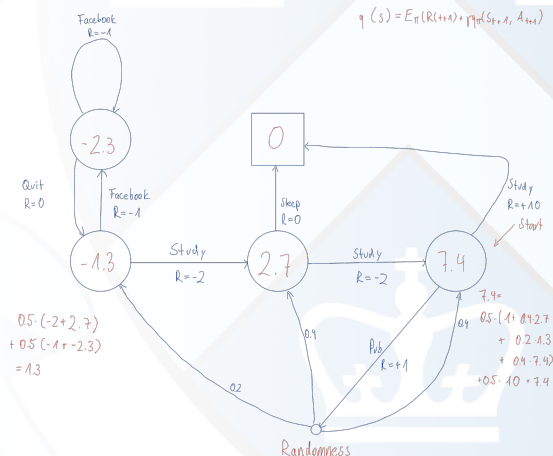
Figure: Markov Decision Process (Credits=David Silver, DeepMind)



# Markov Decision Process

Visualized


Figure: Markov Decision Process (Credits=David Silver, DeepMind)



# Markov Decision Process

Visualized

Figure: Markov Decision Process (Credits=David Silver, DeepMind)



# Definition

... this Bellman-Optimality equation where we take  $\max q_*$  can be solved through Q-learning (which is off-policy).

## Definition

For any finite Markov decision process (FMDP), Q-learning finds an optimal policy in the sense of maximizing the expected value of the total reward over any and all successive steps, starting from the current state. Q-learning can identify an optimal action-selection policy for any given FMDP, given infinite exploration time and a partly-random policy.

## Q-learning formula

$$Q^{new}(s_t, a_t) \leftarrow \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \underbrace{\left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} - \underbrace{Q(s_t, a_t)}_{\text{old value}} \right)}_{\text{new value (temporal difference target)}} \quad (16)$$

temporal difference

# Epsilon-Greedy strategy

At each time step, our agent can either **explore** or **exploit**.

## Carl Example

At lunchtime, Carl can either go exploit (go to a place where he knows the food is good, he will get a great reward) or explore (go to a new place with the goal of finding a place with even higher reward)

# Epsilon-Greedy strategy

Hence, we start with an  $\epsilon$  (exploration rate, the probability that our agent will explore rather than exploit) of 1 which gradually decreases after each episode.

## Pseudo-Code

```
if(random.generate(0,1) > epsilon) :  
    // exploit, choose action with highest Q-value  
else :  
    // explore, choose a random action and see what happens
```



# Numerical Example

Suppose the reward Matrix  $R$  is given by

state / action	A	B	C	D	E	F
A	-	-	-	-	0	-
B	-	-	-	0	-	100
C	-	-	-	0	-	-
D	-	0	0	-	0	-
E	0	-	-	0	-	100
F	-	0	-	-	0	100

# Numerical Example

And we initialize the Q-table...

	A	B	C	D	E	F
A	0	0	0	0	0	0
B	0	0	0	0	0	0
C	0	0	0	0	0	0
D	0	0	0	0	0	0
E	0	0	0	0	0	0
F	0	0	0	0	0	0

... where F is the goal state.

# Numerical Example

Now we are starting in state B and randomly go to F:  
We can update the Q-table in the following way:

## Update Q-values

$$Q^{new}(B, F) \leftarrow \underbrace{Q(B, F)}_{\text{old value}} + \underbrace{0.1}_{\text{learning rate}} \cdot \left( \underbrace{R(B, F)}_{\text{reward}} + \underbrace{0.8}_{\text{discount factor}} \cdot \underbrace{\max_a(Q(F, B), Q(F, E), Q(F, F))}_{\text{estimate of optimal future value}} - \underbrace{Q(B, F)}_{\text{old value}} \right) \quad (17)$$

temporal difference  
new value (temporal difference target)

$$Q^{new}(B, F) \leftarrow \underbrace{0}_{\text{old value}} + \underbrace{0.1}_{\text{learning rate}} \cdot \left( \underbrace{100}_{\text{reward}} + \underbrace{0.8}_{\text{discount factor}} \cdot \underbrace{\max_a(0, 0, 0)}_{\text{estimate of optimal future value}} - \underbrace{0}_{\text{old value}} \right) \quad (18)$$

temporal difference  
new value (temporal difference target)

$$Q^{new}(B, F) = 0 + 0.1 \cdot (100 + 0.8 \cdot 0 - 0) = 10 \quad (19)$$

# Numerical Example

And we update our Q-table

	A	B	C	D	E	F
A	0	0	0	0	0	0
B	0	0	0	0	0	10
C	0	0	0	0	0	0
D	0	0	0	0	0	0
E	0	0	0	0	0	0
F	0	0	0	0	0	0

... and the episode is finished, since F is the goal state.

# Deep reinforcement learning motivation

Problems with vanilla Q-learning:

1. Storing the entries  $Q(s,a)$  for all possible States and actions  $A$  (in our matrix) is not always feasible:

In classical board games:  $|S| = b^d$

where  $b$  is the mean number of allowed moves for a given board position and  $d$  is the depth, the typical number of moves per game.

Chess:  $b = 35$ ,  $d = 80$

2. Transition probability in MDP's is often unknown or can only be approximated

→ Approximate the Q-table with a Neural Net

# Policies, Targets and TD's

- ▶ Looking back at the formula for the Q-Value: We can see we have a Q-Value for  $Q(s,a)$  at time  $t$ , but then we can recalculate this value at  $t+1$  with new information.
- ▶ This value at  $t+1$  (after the action was taken) - the original value is the temporal difference observed in the formula.
- ▶ Naturally, the closer this value gets to zero, the closer we are in converging to the true value.

# Policies, Targets and TD's

We can formalize this problem through the MSE:

$$L(\Theta) = \mathbf{E}_a \|Y - Q(S, A, \Theta)\|_2^2 \quad (20)$$

where  $y_i$  is the temporal difference and  $\theta_i$  are the weights and biases of the network.

# Policies, Targets and TD's

And in order to not play cat and mouse, we introduce a policy and a target network:




Figure: Policy and target networks in DQN's



# Experience replay

- ▶ In order to get better at estimate of the target Q-values we store our trajectory  $(S,A,R,S')$  at each step.
- ▶ We then randomly sample from this replay and train the model based on these results.

# DQN Trading Bot

See uploaded Jupyter Notebook



Sutton, Richard S. and Andrew G. Barto. *Reinforcement learning: an introduction*. Adaptive computation and machine learning. Cambridge, Mass: MIT Press, 1998. 322 pp. ISBN: 9780262193986.



White, Chelsea C. “A survey of solution techniques for the partially observed Markov decision process”. In: *Annals of Operations Research* 32.1 (Dec. 1, 1991), pp. 215–230. ISSN: 1572-9338. DOI: 10.1007/BF02204836. URL: <https://doi.org/10.1007/BF02204836> (visited on 03/13/2022).



# Reinforcement Learning

a practical approach

Prof. Renzo Silva, Konstantin Kuchenmeister, Elijah Hall

*Department of Mathematics*  
*Columbia University*

Mathematics Building  
MATH203

Saturday 9<sup>th</sup> April, 2022